

a single function version of `makeVector` and `cachemean` from the JHU Data Science Specialization on Coursera R Programming course

Alan E. Berger January 24, 2025

available at https://github.com/AlanBerger/R_Programming_GitHub_files

Motivation

The R Programming course in the Johns Hopkins University Data Science Specialization on Coursera (a 10 course sequence) covers many topics in the R Programming language, including *lexical scoping*. As part of a programming exercise involving lexical scoping, the pair of R functions, `makeVector` and `cachemean`, are made available to the class; they are publicly available within the **README.md** file in Professor Roger D. Peng's GitHub site: <https://github.com/rdpeng/ProgrammingAssignment2> A detailed description of how `makeVector` and `cachemean` work is available in an article by Leonard Greski at: <https://github.com/lgreski/datasciencecontent/blob/master/markdown/rprog-breakingDownMakeVector.md>

If you download a copy of this pdf file into your computer, then the web links in it will work.

A call to the `makeVector` function returns an R object that is a list of 4 functions, and also is a *function closure*, well described in Leonard Greski's article, but not referred to by that name. Function closures are also discussed in, for example, https://github.com/AlanBerger/R_Programming_GitHub_files/blob/master/reply-to-where-function-closure-variables-are-8-Jan-2020.md

A good observation in the Discussion Forum in the JHU R Programming class noted that someone could “misuse” the `setmean` function provided by `makeVector`, replacing the current saved (*cached*) value of the mean of a given vector with an incorrect value.

The programming exercise involving `makeVector` and `cachemean` was focused on understanding lexical scoping, and not concerned at that point with producing “library quality” code that would “shield” the user from misadventure. However, I thought it an interesting exercise to write a single function that would retain the primary functionality of `makeVector` and `cachemean`, illustrate use of lexical scoping and the properties of a function closure, be simpler to use, and perhaps be a more concise demonstration of key concepts.

This function, `get_and_cache_mean`, is provided below, along with sample runs. It is heavily commented, and uses long (one might say overly long) variable names to make it easier to follow.

The `get_and_cache_mean` function and example runs

The `get_and_cache_mean` function has no argument. It codes for and returns the function `get_the_mean_of_v`

Since `get_and_cache_mean` defines and returns the function `get_the_mean_of_v`, `get_the_mean_of_v` is also a **function closure** and has the capability of “retaining” variables in its parent environment, and accessing and modifying them (see the referenced articles above, or other sources on function closure). When `get_the_mean_of_v` is called, its argument should be a non-empty numeric vector, call it `v`. `get_the_mean_of_v` calculates and returns the mean of `v`, and “stores” the calculated mean. The key point of this programming example is to use lexical scoping and the properties of a function closure so that if `get_the_mean_of_v` has been called before in the current R session, and the most recent call was with the same value of `v` as the present call, then `get_the_mean_of_v` will “fetch” the previously calculated value of the mean that was stored in the parent environment of `get_the_mean_of_v` (possible since `get_the_mean_of_v` is a function closure), rather than recalculating the mean.

If want to copy code, copy it from the Rmd (R Markdown) file, not the pdf file (copying code from a pdf file in GitHub might entrain formatting characters unacceptable to R).

```
get_and_cache_mean <- function() {  
  # call it via:      get_the_mean_of_v <- get_and_cache_mean()  
  
  # get_and_cache_mean returns the function get_the_mean_of_v  
  
  # get_the_mean_of_v will have 1 argument, v (a numeric vector),  
  # and will calculate the mean of v,  
  # and get_the_mean_of_v will also be a function closure  
  
  m <- NULL # initialize the mean m of v to be NULL so  
  #          know need to calculate it  
  
  v_for_which_calculated_the_mean <- NULL  
  # the v for which have calculated the mean, initially NULL  
  
  # define the function to be returned by get_and_cache_mean  
  
  get_the_mean_of_v <- function(v = numeric()) {  
    # calculate the mean of v, or when appropriate fetch and  
    # return the cached (saved) mean  
  
    # do some testing for whether the value that get_the_mean_of_v  
    # was called with is "admissible" for what get_the_mean_of_v  
    # is designed to do  
  
    if(is.numeric(v) == FALSE)  
      stop("the argument v of get_the_mean_of_v is not numeric")  
    if(is.vector(v) == FALSE)
```

```

    stop("the argument v of get_the_mean_of_v is not a vector")
    if(length(v) < 1) stop("length of v is < 1")
#   note a variable that is a single value is, in R,
#   a vector of length 1
#   a matrix and a data frame are examples of R objects
#   that are not vectors

    m_from_parent_envir <- m
#   use R's lexical scoping to acquire m

    v_from_parent_environment <- v_for_which_calculated_the_mean
#   use R's lexical scoping to acquire
#   v_for_which_calculated_the_mean

#   test whether there is a cached (saved) mean that is
#   appropriate to use

#   need to be VERY careful with any NULL values
#   use the is.null function to test for a potential NULL value;
#   if have verified that
#   the value is not NULL,
#   then can do regular logical tests or arithmetic
#   on it (the length of the NULL R object is 0)
#
if(!is.null(m_from_parent_envir) && !is.null(v_from_parent_environment)) {
#   if got to here, neither is NULL, so have cached a mean
#
#   test whether the value of v used in the call to get_the_mean_of_v
#   is the one for which the cached mean has been calculated

    if(identical(v_from_parent_environment, v)) {
#       note for example c(1, 2, 3) is numeric while 1:3 is
#       the vector of integers c(1L, 2L, 3L) so they will not be
#       identical R objects
#       (even though their entries are equal numbers)

#       if pass the if test above, can use the cached (saved) mean

        message("getting cached value of the mean of v")
        return(m_from_parent_envir) # returns the cached mean
#       and exits the get_the_mean_of_v function

    } # closing brace for the if test immediately above
} # closing brace for the if test above that checks for NULL values

```

```

#      if got to here, need to calculate the mean of v
#      which is the v that is the argument of the current call
#      of get_the_mean_of_v

      v_for_which_calculated_the_mean <- v
#          use the superassignment operator <- to
#          update v_for_which_calculated_the_mean
#          in the parent environment of get_the_mean_of_v

#          the parent environment is a preserved version of
#          the "execution time" environment of
#          get_and_cache_mean

#          The R statement
#          get_the_mean_of_v <- get_and_cache_mean()
#          defines and returns the function get_the_mean_of_v
#          and hence get_the_mean_of_v
#          is also a "function closure"

#          As a crucial consequence of get_the_mean_of_v being a
#          function closure, the variables (R objects) m
#          and v_for_which_calculated_the_mean
#          defined in the original call to get_and_cache_mean
#          are preserved and available within the
#          parent environment of get_the_mean_of_v
#          so available, if properly accessed, to
#          the function get_the_mean_of_v

      mean_of_v <- mean(v) # note if v were v = numeric(),
#                          i.e., a numeric vector of length 0,
#                          its mean would be NaN
      m <- mean_of_v # use the superassignment operator <-
#                  to cache the mean of v in the parent
#                  environment of get_the_mean_of_v

      message("returning the calculated mean of v, NOT a cached value")
      message("have now cached the mean of the current value of v")
      return(mean_of_v)
} # ending brace for the function get_the_mean_of_v

} # ending brace for the function get_and_cache_mean

#####

```

```
# Now do some sample runs
```

```
get_the_mean_of_v <- get_and_cache_mean()
```

```
# get_the_mean_of_v is a function and also is a function closure
```

```
v <- c(6, 7, 8) # simple example
```

```
get_the_mean_of_v(v)
```

```
## returning the calculated mean of v, NOT a cached value
```

```
## have now cached the mean of the current value of v
```

```
## [1] 7
```

```
# check that the mean was cached in the previous
```

```
# call of get_the_mean_of_v
```

```
get_the_mean_of_v(v)
```

```
## getting cached value of the mean of v
```

```
## [1] 7
```

```
new_v <- c(1, 2, 3)
```

```
get_the_mean_of_v(new_v) # check works with new vector
```

```
## returning the calculated mean of v, NOT a cached value
```

```
## have now cached the mean of the current value of v
```

```
## [1] 2
```

```
get_the_mean_of_v(new_v) # again, check cached the mean
```

```
## getting cached value of the mean of v
```

```
## [1] 2
```

```
# note c(1, 2, 3) is numeric while 1:3 is
```

```
# the vector of integers c(1L, 2L, 3L) so they will not be
```

```
# identical R objects, and so
```

```
# get_the_mean_of_v will recalculate the mean in this case:
```

```
get_the_mean_of_v(1:3)
```

```
## returning the calculated mean of v, NOT a cached value
## have now cached the mean of the current value of v
```

```
## [1] 2
```

```
get_the_mean_of_v(1:3) # again, check cached the mean
```

```
## getting cached value of the mean of v
```

```
## [1] 2
```

```
## the tests below will give errors and put the R session into Browse
## (note they stop R Markdown, so not run)
## exit Browse by hitting the Escape key
```

```
## get_the_mean_of_v("character string")
# check non-numeric argument gets error
```

```
## matr <- matrix(1:12, 3, 4) # 3 row by 4 column matrix
## get_the_mean_of_v(matr)
# check argument that is not a vector gets error
```

```
## get_the_mean_of_v(numeric(0))
# check argument that has length < 1 gets error
```

Hope this is helpful in understanding R's lexical scoping and how a function closure works. Note the corresponding version of the code above for the inverse of a matrix is NOT what is required for the second assignment in the R Programming course.

=====

This work is licensed under the Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License. To view a copy of this license, visit <https://creativecommons.org/licenses/by-nc-sa/4.0/> or send a letter to Creative Commons, PO Box 1866, Mountain View, CA 94042, USA. There is a full version of this license at this web site: <https://creativecommons.org/licenses/by-nc-sa/4.0/legalcode>

The material above was taken/modified from a post of mine in the Discussion Forum for the R Programming Course in the Johns Hopkins Data Science Specialization on Coursera. As such Coursera and Coursera authorized Partners retain additional rights to that material as described in their "Terms of Use" <https://www.coursera.org/about/terms>

Note the reader should not infer any endorsement or recommendation or approval for the material in this article from any of the sources or persons cited above or any other entities mentioned in this article.